

Assignment 4 MTL458

AVL Tree

Implementation:

In our implementation, we have decided on the following approach: We allow as many read operations as concurrently as needed (contains, in order), and we allow only one write operation at a time (insert, delete). Also, whenever we encounter a delete operation, we ensure that all previously encountered inserts are finished.

We have ensured that no deadlocks occur and there are no race conditions. We have also ensured that there is no starvation.

Comparison:

Comparing the behaviour of a concurrent AVL tree implementation with a traditional (non-concurrent) AVL tree is essential to understand the advantages and trade-offs of each approach. We discuss these below:

1. Non-Determinant Behaviour

Due to concurrency in operations, the structure of the AVL Tree might change drastically, and the in-order and pre-order traversals might be vastly different. This is expected but sometimes not encouraged.

2. Performance Metrics:

	Concurrent AVL Tree	Traditional AVL Tree
Insertion performance	may experience increased overhead due to synchronization, especially in high-contention scenarios. Performance may be affected by the need for locks or atomic operations.	doesn't have the added overhead of synchronization. It typically performs well in insertion operations.
Deletion performance	may experience additional overhead during deletion. The synchronization mechanisms can impact performance.	may perform deletion operations more efficiently, as it doesn't have to deal with synchronization concerns.
Search(contains) performance	Both implementations are likely to have similar performance in search operations, as no write operations are involved, and synchronization overhead is minimal.	
In-order Traversal performance	the performance may vary depending on the synchronization mechanisms used. Contention among threads may lead to suboptimal performance.	can typically perform in-order traversal efficiently as it doesn't have to deal with concurrency-related challenges.
Overall performance	Multiple concurrent reads might result in a good overall performance, but overhead introduced might counter it.	-

3. Memory usage:

Concurrent AVL Tree	Traditional AVL Tree
Node Overhead: tree may have additional memory overhead per node due to synchronization mechanisms. This overhead can include locks, flags, or other data required to manage concurrent access. The memory used per node will depend on the synchronization method chosen.	Node Overhead: trees have a smaller memory overhead per node because they don't require synchronization mechanisms. They only contain the essential components, such as key, value, and pointers to children.
Additional Data Structures: might use additional data structures to manage concurrent operations, such as thread-specific data, contention management structures, or atomic variables for synchronization. These data structures can consume extra memory.	Simplicity: simplicity can lead to less memory overhead. There are no additional data structures or atomic variables to manage concurrent access.
Node Payload: stores the same data payload (keys, values, etc.) as the traditional AVL tree, so the memory usage related to the payload itself will be the same.	

Examples:

1.

Input	Output	Performance
insert 3 insert 4 insert 2 delete 2 contains 3 in order exit	yes 3 4 3 4	Time with concurrency: 0.015 s Time without concurrency: 0.016s

2.

Input	Output	Performance
insert 5	3 5 6 10 14 15 23 24 25 30	Time with concurrency: 0.020 s
insert 10	no	
insert 15	3 5 6 10 14 15 24 25 30	Time without concurrency: 0.018s
insert 25	3 5 6 10 14 15 23 24 25 30 37	
insert 30	no	
insert 14	3 5 6 10 14 15 23 24 25 30 36	
insert 23	15 10 5 3 6 14 25 24 23 36 30	
insert 6		
insert 24		
insert 30		
insert 3		
in order		
delete 23		
contains 23		
in order		
insert 23		
delete 36		
insert 37		
in order		
insert 36		
delete 37		
contains 37		
in order		
exit		

Thus, we can see that there is marginal difference in this case for AVL Trees implemented using concurrency, limited to testing conditions.

Suggested Improvements: We might consider implementing locks per node rather than a lock for the whole tree. However, as always, it also adds further complexity at the cost of concurrency.